

OPENMP/TAU - Performance Analysis

Contents

1	Introduction	2
1.1	Modification of code for parallelization	2
2	Results	3
2.1	Time consumption	3
2.1.1	Serial	3
2.1.2	2 threads	3
2.1.3	4 threads	4
2.1.4	8 threads	4
2.2	Problem size	5
2.3	Level 3 Cache Misses	6
2.3.1	2 threads	6
2.3.2	4 threads	7
2.3.3	8 threads	8
2.4	Total Instructions	9
2.4.1	2 threads	9
2.4.2	8 threads	10
3	Ofast	13
4	Conclusion	14

1 Introduction

1.1 Modification of code for parallelization

We use the following OPENMP commands to parallelize our code. These especially refers to the double for-loop performing the `laplace_step`, that is, the most time-consuming part of all the program.

```
...

float laplace_step(float *in, float *out, int n)
{
    int i, j;
    float error=0.0f;
    #pragma omp for nowait
    for ( j=1; j < n-1; j++ )
        for ( i=1; i < n-1; i++ )
        {
            ...
        }
    return error;
}
...

int main(int argc, char** argv)
{
    ...
    omp_set_num_threads(2); //4,8
    int iter = 0;
    #pragma omp parallel default(none) firstprivate(iter,temp,A,error)
        shared(n, iter_max)
    {
        while (error > tol*tol && iter < iter_max )
        {
            ...
        }

        #pragma omp master
        {
            error = sqrtf( error );
            printf("Total Iterations: %5d, ERROR: %0.6f, ", iter, error);
            printf("A[%d][%d]= %0.6f\n", n/128, n/128, A[(n/128)*n+n/128]);
        }
    }
    ...
}
```

In the `int main`, we introduce a OMP section with declaring private the parameters `iter`, `temp`, `A` and `error`, since each thread handles own copies of these variables. On the other hand, the constant global variables `n` and `iter_max` can be shared among all the threads. In the `laplace_step` function, we use the `pragma_omp_for_nowait` command, which branches the process to perform the double for-loop.

2 Results

2.1 Time consumption

2.1.1 Serial

We use the original code for the Laplace application provided and get the results of the `perf stat` for serial execution with the `gcc` compiler and with a mesh of $n = 1000$:

Total Iterations: 1000, ERROR: 0.015535, A[7][7]= 0.016615

Performance counter stats for './lap_2 1000':

31250.063048	task-clock (msec)	#	0.998 CPUs utilized
104,364,169,322	cycles	#	3.340 GHz
115,660,649,729	instructions	#	1.11 insn per cycle

31.308894954 seconds time elapsed

For all the following, we execute the code with the `tau_cc.sh` compiler and a mesh of $n = 1000$. The results are obtained from the `perf` tool and the TAU `pprof` command.

2.1.2 2 threads

Jacobi relaxation Calculation: 1000 x 1000 mesh, maximum of 1000 iterations

Total Iterations: 1000, ERROR: 0.015535, A[7][7]= 0.016615

Performance counter stats for './lap_2_2 1000':

112963,788825	task-clock (msec)	#	1,987 CPUs utilized
371.475.376.651	cycles	#	3,288 GHz
706.115.703.767	instructions	#	1,90 insns per cycle

56,845403345 seconds time elapsed

FUNCTION SUMMARY (mean):

%Time	Exclusive msec	Inclusive total msec	#Call	#Subrs	Inclusive Name usec/call
100.0	7	56,826	1	1	56826868 .TAU application
100.0	0.798	56,817	1	1001.5	56817807 parallel (parallel begin/end)
99.7	1	56,684	1000	1000	56684 laplace_step
99.7	56,612	56,683	1000	199944	56683 for (loop body)
50.0	0.286	28,410	0.5	1.5	56821454 main
50.0	0.159	28,409	0.5	0.5	56818177 parallel (parallel fork/join)
0.2	132	132	1	0	132790 parallel (barrier enter/exit)
0.1	47	47	99973	0	0 max_error [THROTTLED]
0.0	23	23	99971.5	0	0 stencil [THROTTLED]
0.0	1	1	1	0	1352 laplace_init
0.0	0.0205	0.0205	0.5	0	41 master (master begin/end)

The execution times increases when we use the `tau_cc.sh` compiler. Remarkable is the fact, that it can be seen a strong increase of cycles and instructions (approx. by the factor 7).

2.1.3 4 threads

Jacobi relaxation Calculation: 1000 x 1000 mesh, maximum of 1000 iterations
Total Iterations: 1000, ERROR: 0.012964, A[7][7]= 0.016615

Performance counter stats for './lap_2_tau 1000':

116704,631422	task-clock (msec)	#	3,957 CPUs utilized
372.373.716.617	cycles	#	3,191 GHz
706.643.925.605	instructions	#	1,90 insns per cycle

29,494571057 seconds time elapsed

%Time	Exclusive msec	Inclusive total msec	#Call	#Subrs	Inclusive Name usec/call
100.0	3	29,366	1	1	29366954 .TAU application
100.0	0.49	29,362	1	1001.25	29362354 parallel (parallel begin/end)
99.4	0.812	29,187	1000	1000	29188 laplace_step
99.4	29,133	29,186	1000	197776	29187 for (loop body)
25.0	0.172	7,341	0.25	0.75	29366667 main
25.0	0.118	7,340	0.25	0.25	29362843 parallel (parallel fork/join)
0.6	174	174	1	0	174124 parallel (barrier enter/exit)
0.1	31	31	98889	0	0 max_error [THROTTLED]
0.1	22	22	98887.2	0	0 stencil [THROTTLED]
0.0	0.784	0.784	0.5	0	1568 laplace_init
0.0	0.0138	0.0138	0.25	0	55 master (master begin/end)

Here we see a reduction of the time by the half. The major part of 74,9% is taken by the double for-loop inside the `laplace_step` function. The cycles and instructions stay at the same level as before.

2.1.4 8 threads

For the case of 8 threads, we will take a closer view at the behaviour under change of problem size. As before, we start with a mesh size of $n = 1000$.

We compiled and executed the program with 8 threads via the remote access. In difference to the execution in the lab, we only have the option to use 4 CPUs. So it is reasonable that we obtain similar results (~ 29 sec) as in the case of 4 threads.

We expect that the execution time on 8 CPUs would approximately decrease by half.

FUNCTION SUMMARY (mean):

%Time	Exclusive msec	Inclusive total msec	#Call	#Subrs	Inclusive Name usec/call
100.0	2	29,188	1	1	29188116 .TAU application
100.0	0.428	29,184	1	1001.12	29184637 parallel (parallel begin/end)
99.8	0.659	29,119	1000	1000	29119 laplace_step
99.8	29,063	29,118	1000	96512.1	29118 for (loop body)
12.5	0.078	3,649	0.125	0.375	29194622 main
12.5	0.0839	3,648	0.125	0.125	29191336 parallel (parallel fork/join)
0.2	65	65	1	0	65145 parallel (barrier enter/exit)

0.1	29	29	48255.8	0	1 stencil [THROTTLED]
0.1	25	25	48256.4	0	1 max_error [THROTTLED]
0.0	0.333	0.333	0.25	0	1331 laplace_init
0.0	0.00762	0.00762	0.125	0	61 master (master begin/end)

2.2 Problem size

The effect of the change of problem size on the performance of the program was monitored and the results are presented below. The program was compiled with 8 threads for this section.

Jacobi relaxation Calculation: 1000 x 1000 mesh, maximum of 1000 iterations
Total Iterations: 1000, ERROR: 0.010200, A[7][7]= 0.016615

Performance counter stats:

116659,711407	task-clock (msec)	#	3,975 CPUs utilized
372.143.041.826	cycles	#	3,190 GHz
706.639.778.898	instructions	#	1,90 insns per cycle

29,351555976 seconds time elapsed

Jacobi relaxation Calculation: 500 x 500 mesh, maximum of 1000 iterations
Total Iterations: 1000, ERROR: 0.009603, A[3][3]= 0.016854

Performance counter stats:

29402,502412	task-clock (msec)	#	3,712 CPUs utilized
93.508.818.896	cycles	#	3,180 GHz
176.784.479.052	instructions	#	1,89 insns per cycle

7,921265885 seconds time elapsed

Jacobi relaxation Calculation: 100 x 100 mesh, maximum of 1000 iterations
Total Iterations: 1000, ERROR: 0.013125, A[0][0]= 1.000000

Performance counter stats:

1488,872256	task-clock (msec)	#	2,258 CPUs utilized
4.360.606.365	cycles	#	2,929 GHz
7.885.220.503	instructions	#	1,81 insns per cycle

0,659373830 seconds time elapsed

As expected, the number of cycles, instructions as well as the execution time decrease approximately with the square of the change of the problem size: Comparing the stats of $n = 1000$ and $n = 500$, all the parameters are 4 times as small in the latter case. Comparing $n = 500$ and $n = 100$, we obtain the same behaviour, except the time, which is more than expected. We suppose that reason for this is the overall short execution times below 1 second, which noises the time result. Especially there are processes like the establishment of the parallel region, that need to be done regardless which problem size is used.

2.3 Level 3 Cache Misses

For Level 3 (L3) Cache misses, we measure how many times an access call was made to the L3 cache and the data requested was not found, meaning it had to be fetched from the main memory, normally resulting in an increase latency. With MULTI_PAPI_L3_TCM as the activated PAPI_METRIC, we can obtain measurements of the Level 3 cache misses. Again we compare results in the case of 2,4 and 8 threads with a mesh size of $n = 1000$.

2.3.1 2 threads

NODE 0;CONTEXT 0;THREAD 0:

%Time	Exclusive counts	Inclusive total counts	#Call	#Subrs	Count/Call	Name
100.0	5455	8102	1	1	8102	.TAU application
32.7	1190	2647	1	3	2647	main
16.4	622	1328	1	1	1328	parallel (parallel fork/join)
8.7	118	706	1	1002	706	parallel (parallel begin/end)
4.3	42	352	1000	1000	0	laplace_step
3.8	280	310	1000	200002	0	for (loop body)
2.9	234	234	1	0	234	master (master begin/end)
1.6	129	129	2	0	64	laplace_init
0.2	19	19	100001	0	0	stencil [THROTTLED]
0.1	11	11	100001	0	0	max_error [THROTTLED]
0.0	2	2	1	0	2	parallel (barrier enter/exit)

NODE 0;CONTEXT 0;THREAD 1:

%Time	Exclusive counts	Inclusive total counts	#Call	#Subrs	Count/Call	Name
100.0	464	890	1	1001	890	parallel (parallel begin/end)
95.8	-37	853	1	1	853	.TAU application
25.8	230	230	1	0	230	parallel (barrier enter/exit)
22.0	30	196	1000	1000	0	laplace_step
18.7	150	166	1000	198046	0	for (loop body)
1.0	9	9	99022	0	0	stencil [THROTTLED]
0.8	7	7	99024	0	0	max_error [THROTTLED]

FUNCTION SUMMARY (mean):

%Time	Exclusive counts	Inclusive total counts	#Call	#Subrs	Count/Call	Name
99.6	2709	4478	1	1	4478	.TAU application
29.4	595	1324	0.5	1.5	2647	main
17.7	291	798	1	1001.5	798	parallel (parallel begin/end)
14.8	311	664	0.5	0.5	1328	parallel (parallel fork/join)
6.1	36	274	1000	1000	0	laplace_step
5.3	215	238	1000	199024	0	for (loop body)
2.6	117	117	0.5	0	234	master (master begin/end)
2.6	116	116	1	0	116	parallel (barrier enter/exit)
1.4	64.5	64.5	1	0	64	laplace_init
0.3	14	14	99511.5	0	0	stencil [THROTTLED]




Figure 1: paraprof output for the Level 3 total cache misses count



Figure 2: paraprof legend of colours representing the different program parts




Figure 3: The paraprof output shows that the majority of cache misses happen during the for-loop, a small amount at the rest of the laplace_step

For the number of L3 cache misses we expected to obtain the highest numbers during the double loop in the laplace_step. However, viewing the pprof results, we can see that in both threads, the TAU application itself consumes the majority of cache misses (67.3% in thread 0, 70% in thread 1). In contrast to this, the for-loop, which was assumed to be most time consumptive, is only responsible for a small fraction of cache misses (1.1% and 17.7%). Seeing the output of paraprof, our expectations are confirmed.

2.3.2 4 threads

FUNCTION SUMMARY (mean):

%Time	Exclusive counts	Inclusive total counts	#Call	#Subrs	Count/Call	Name
100.0	1387	1.302E+06	1	1	1301916	.TAU application
99.8	7316	1.3E+06	1	1001.25	1299769	parallel (parallel begin/end)

99.3	2.422E+04	1.292E+06	1000	1000	1292	laplace_step
97.4	1.268E+06	1.268E+06	1000	192937	1268	for (loop body)
25.1	460.2	3.263E+05	0.25	0.75	1305293	main
25.0	191.5	3.258E+05	0.25	0.25	1303019	parallel (parallel fork/join)
0.0	108.2	108.2	0.5	0	216	laplace_init
0.0	100	100	1	0	100	parallel (barrier enter/exit)
0.0	98.25	98.25	0.25	0	393	master (master begin/end)
0.0	7	7	96468	0	0	stencil [THROTTLED]
0.0	6.5	6.5	96469.2	0	0	max_error [THROTTLED]

Metric: PAPL3_TCM
Value: Exclusive





Figure 4: exclusive level 3 cache misses count

2.3.3 8 threads

FUNCTION SUMMARY (mean):

%Time	Exclusive counts	Inclusive total counts	#Call	#Subrs	Count/Call	Name
100.0	752.2	6.267E+05	1	1	626672	.TAU application
99.8	4846	6.255E+05	1	1001.12	625518	parallel (parallel begin/end)
99.0	1.614E+04	6.206E+05	1000	1000	621	laplace_step
96.4	6.042E+05	6.044E+05	1000	199216	604	for (loop body)
13.7	231.2	8.608E+04	0.125	0.375	688645	main
13.7	111.6	8.579E+04	0.125	0.125	686330	parallel (parallel fork/join)
0.0	191.9	191.9	99608.5	0	0	max_error [THROTTLED]
0.0	67.75	67.75	1	0	68	parallel (barrier enter/exit)
0.0	58.12	58.12	0.25	0	232	laplace_init
0.0	45.12	45.12	0.125	0	361	master (master begin/end)
0.0	4.75	4.75	99607.5	0	0	stencil [THROTTLED]

Executing with 8 threads, we observe that in mean, the total inclusive counts increases from 4500 to 62000. The `parallel fork/join` consumes 13.7% of all cache misses, the double for-loop is responsible for 83.7% of them.



Figure 5: The other parts of the program play an insignificant role compared to the double for-loop, whilst there is some variability between the threads




Figure 6: comparison of 4 and 8 threads cache misses for different problem sizes on a logarithmic scale

2.4 Total Instructions

2.4.1 2 threads

NODE 0;CONTEXT 0;THREAD 0:

%Time	Exclusive counts	Inclusive total counts	#Call	#Subrs	Count/Call	Name
100.0	1.196E+07	3.757E+09	1	1	3756922236	.TAU application
99.7	8.471E+04	3.745E+09	1	3	3744966401	main
99.7	2.261E+05	3.745E+09	1	1	3744841689	parallel (parallel fork/join)
99.7	1.279E+06	3.745E+09	1	1002	3744615637	parallel (parallel begin/end)
99.6	1.802E+06	3.743E+09	1000	1000	3743327	laplace_step
99.6	3.583E+09	3.742E+09	1000	200002	3741525	for (loop body)
2.1	7.95E+07	7.95E+07	100001	0	795	max_error [THROTTLED]
2.1	7.92E+07	7.92E+07	100001	0	792	stencil [THROTTLED]
0.0	4E+04	4E+04	2	0	20001	laplace_init
0.0	8133	8133	1	0	8133	master (master begin/end)
0.0	815	815	1	0	815	parallel (barrier enter/exit)

NODE 0;CONTEXT 0;THREAD 1:

%Time	Exclusive counts	Inclusive total counts	#Call	#Subrs	Count/Call	Name
100.0	5.672E+04	3.754E+09	1	1	3753769193	.TAU application
100.0	1.315E+06	3.754E+09	1	1001	3753712468	parallel (parallel begin/end)
99.9	1.843E+06	3.75E+09	1000	1000	3750294	laplace_step
99.9	3.59E+09	3.748E+09	1000	198046	3748451	for (loop body)
2.1	8.001E+07	8.001E+07	99024	0	808	max_error [THROTTLED]
2.1	7.843E+07	7.843E+07	99022	0	792	stencil [THROTTLED]
0.1	2.103E+06	2.103E+06	1	0	2102895	parallel (barrier enter/exit)

FUNCTION SUMMARY (mean):

%Time	Exclusive counts	Inclusive total counts	#Call	#Subrs	Count/Call	Name
100.0	6.006E+06	3.755E+09	1	1	3755345714	.TAU application
99.8	1.297E+06	3.749E+09	1	1001.5	3749164052	parallel (parallel begin/end)
99.8	1.823E+06	3.747E+09	1000	1000	3746811	laplace_step
99.7	3.586E+09	3.745E+09	1000	199024	3744988	for (loop body)
49.9	4.236E+04	1.872E+09	0.5	1.5	3744966401	main
49.9	1.13E+05	1.872E+09	0.5	0.5	3744841689	parallel (parallel fork/join)
2.1	7.976E+07	7.976E+07	99512.5	0	801	max_error [THROTTLED]
2.1	7.881E+07	7.881E+07	99511.5	0	792	stencil [THROTTLED]
0.0	1.052E+06	1.052E+06	1	0	1051855	parallel (barrier enter/exit)
0.0	2E+04	2E+04	1	0	20001	laplace_init
0.0	4066	4066	0.5	0	8133	master (master begin/end)

Here, the pprof output draws a clear picture: $\sim 97\%$ of the execution time and $\sim 96\%$ of the instructions are consumed during the nested for-loop, marginalizing the instructions for TAU and parallel.




Figure 7: The main (orange) as well as laplace_init (blue) only happens in the main thread, which is reflected in this normalized paraprof output

2.4.2 8 threads

Problem size $n = 1000$

FUNCTION SUMMARY (mean):

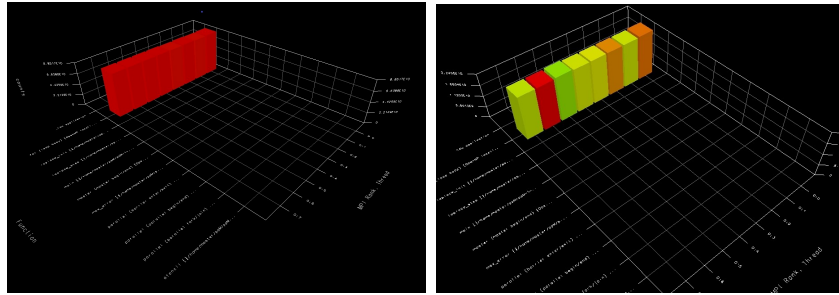
%Time	Exclusive counts	Inclusive total counts	#Call	#Subrs	Count/Call	Name
100.0	1.663E+06	8.85E+10	1	1	88503563932	.TAU application
100.0	1.301E+06	8.85E+10	1	1001.12	88501648242	parallel (parallel begin/end)
100.0	1.958E+06	8.85E+10	1000	1000	88498507	laplace_step
100.0	8.833E+10	8.85E+10	1000	199216	88496549	for (loop body)
12.5	1.051E+04	1.108E+10	0.125	0.375	88672931312	main

12.5	9.826E+04	1.108E+10	0.125	0.125	88671696272	parallel (parallel fork/join)
0.1	9.027E+07	9.027E+07	99608.5	0	906	max_error [THROTTLED]
0.1	7.889E+07	7.889E+07	99607.5	0	792	stencil [THROTTLED]
0.0	1.839E+06	1.839E+06	1	0	1838586	parallel (barrier enter/exit)
0.0	1.439E+05	1.439E+05	0.25	0	575460	laplace_init
0.0	1052	1052	0.125	0	8418	master (master begin/end)


Metric: PAPI_TOT_INS
Value: Exclusive



Figure 8: Compared to the Cache misses, the number of instructions are nearly constant between the 8 threads, which leads to a very small standard deviation



(a) Total instructions with 8 threads and $n = 1000$ (b) Total instructions with 8 threads and $n = 500$



(c) Total instructions with 8 threads and $n = 100$

Problem size $n = 500$

FUNCTION SUMMARY (mean):

%Time	Exclusive counts	Inclusive counts	#Call	#Subrs	Count/Call	Name
100.0	1.58E+06	2.23E+10	1	1	22302478701	.TAU application

100.0	1.301E+06	2.23E+10	1	1001.12	22300556447	parallel (parallel begin/end)
100.0	1.952E+06	2.23E+10	1000	1000	22297415	laplace_step
100.0	2.213E+10	2.23E+10	1000	197794	22295463	for (loop body)
12.6	1.051E+04	2.82E+09	0.125	0.375	22557877852	main
12.6	2.873E+05	2.82E+09	0.125	0.125	22557434948	parallel (parallel fork/join)
0.4	8.985E+07	8.985E+07	98898.1	0	909	max_error [THROTTLED]
0.4	7.833E+07	8.278E+07	98896.9	0.5	837	stencil [THROTTLED]
0.0	1.839E+06	1.839E+06	1	0	1838587	parallel (barrier enter/exit)
0.0	4.485E+04	4.485E+04	0.25	0	179404	laplace_init
0.0	1050	1050	0.125	0	8396	master (master begin/end)

Problem size $n = 100$

FUNCTION SUMMARY (mean):

%Time	Exclusive counts	Inclusive total counts	#Call	#Subrs	Count/Call	Name
100.0	1.593E+06	1.201E+09	1	1	1201193798	.TAU application
99.8	1.301E+06	1.199E+09	1	1001.12	1199314411	parallel (parallel begin/end)
99.6	1.9E+06	1.196E+09	1000	1000	1196173	laplace_step
99.4	1.038E+09	1.194E+09	1000	196502	1194274	for (loop body)
13.0	1.053E+04	1.556E+08	0.125	0.375	1244577237	main
13.0	2.704E+05	1.556E+08	0.125	0.125	1244452974	parallel (parallel fork/join)
6.5	7.831E+07	7.831E+07	98251.8	0	797	max_error [THROTTLED]
6.5	7.781E+07	7.781E+07	98250.6	0	792	stencil [THROTTLED]
0.2	1.839E+06	1.839E+06	1	0	1838592	parallel (barrier enter/exit)
0.0	5000	5000	0.25	0	20001	laplace_init
0.0	1022	1022	0.125	0	8179	master (master begin/end)




Figure 10: comparison of 4 and 8 threads total instructions for different problem sizes

3 Ofast

Ofast with gcc, 8 threads

Jacobi relaxation Calculation: 1000 x 1000 mesh, maximum of 1000 iterations
Total Iterations: 1000, ERROR: 0.010200, A[7][7]= 0.016615

Performance counter stats for './lap_2_open_fast 1000':

926,100107	task-clock (msec)	#	2,014 CPUs utilized
2.441.747.945	cycles	#	2,637 GHz
4.105.164.512	instructions	#	1,68 insns per cycle

0,459805784 seconds time elapsed

Ofast with tau_cc.sh, 8 threads

Jacobi relaxation Calculation: 1000 x 1000 mesh, maximum of 1000 iterations
Total Iterations: 1000, ERROR: 0.010200, A[7][7]= 0.016615

Performance counter stats:

100208,549694	task-clock (msec)	#	3,933 CPUs utilized
319.593.850.400	cycles	#	3,189 GHz
602.953.311.419	instructions	#	1,89 insns per cycle

25,476217679 seconds time elapsed




Figure 11: Comparison of execution times with the optimization tool Ofast.

We compiled the program under different scenarios with the compiler flag `-Ofast`. For the serial execution, we obtained a massive time decrease from ~ 31 sec to ~ 1.3 sec. When we added the parallelization with `fopenmp`, execution time was even reduced to 0.45 sec. However, compiling with `tau_cc.sh`, the effect of `Ofast` was only minor. Adding `Ofast` gave a time reduction from ~ 29 sec to ~ 25 sec.

4 Conclusion

As part of this exercise and as presented in this report, we have implemented parallelization of the laplace program using OMP and monitored several metrics and their impact on performance.

Execution time

With respect to the program compiled with `tau` compiler, the execution time decreased as the number of threads was increased. In most cases, this decrease was proportional to the change in number of threads. However, not much difference is noticed between using 4 threads on a 4-core CPU and using 8 threads on the same CPU.

Problem size

The execution time increased quadratically in proportion to the size of the mesh when the mesh size was increased, and this was as expected as the major part of the program has a quadratic complexity ($\mathcal{O}(n^2)$).

Level 3 cache misses

It was noticed that the amount of L3 cache misses increased as the number of threads was increased.

Total instructions

In terms of instructions it was noticed that there was little variation between the number of instructions executed by each thread when the program was parallelized. The total instructions did remain the same even with an increase in the number of threads, however the average did reduce linearly with respect to the number of threads.

Despite the fact that the use of the OFAST compiler flag had a significant impact on the performance of the program when compiled with `gcc`, very little change was noticed when combined with the Tau compiler.